

# Digital Image Processing (CS7.404)

## Assignment 0

Soham Vaishnav | 2022112002

soham.vaishnav@research.iiit.ac.in

August 11, 2024

### Abstract

Assignment 0 deals with fundamental image processing techniques including manipulation of color as well as grayscale images. The tasks range from basic ones like reading from and writing to images to converting grayscale images into color ones using pseudocoloring. Some of the tasks are relatively complex like Chroma Keying, reading from and writing into a video (essentially an array of frames) and creating a video wherein one image transforms into another image via a smooth transition using some special effect like fading or sliding. The assignment entails using majorly three libraries - OpenCV, Numpy and Matplotlib. OS library has also been used for integrating file and directory paths and to validate their existence.

## 0.1 Task 1 - Reading from Image

This task involves reading an image from the folder into an array for any future use. I have used **OpenCV** for the majority of the task.

```
def ImageRead(Image_Dir: str, file_name: str):  
    """  
    Read an image into an array and return the array as well a boolean value to denote if it is  
    a colored image (RGB) or not.  
  
    Arguments:  
    Image_Dir = path to directory containing the images.  
    file_name = name of the image file.  
    """  
    assert os.path.exists(Image_Dir), f"{Image_Dir} does not exist!" #checking for input path validity  
    image_path = os.path.join(Image_Dir, file_name)  
  
    assert os.path.exists(image_path), f"{image_path} does not exist!" #checking for image path validity  
    image = cv2.imread(image_path, cv2.IMREAD_UNCHANGED)  
    if (len(image.shape) == 3 and image.shape[2] == 3):  
        image = Convert(image, is_BGR = True, to_RGB = True, is_RGB = False, to_BGR = False)  
        print("Image " + file_name + " is read into an array and is stored in RGB format"  
              " - 3 channels detected.")  
        return image, True  
    print("Image " + file_name + " is read into an array and is stored in Grayscale format"  
          " - only 1 channel detected.")  
    return image, False
```

Figure 1: Function for reading from an image into an array

The inputs to the function are the image directory and the name of the image we want to read from. The function returns the image array as well as a boolean denoting its type (color or grayscale). Note that the image returned is in RGB format after conversion from BGR which OpenCV returns. The outputs can be seen in the main jupyter book.

## 0.2 Task 2 - Writing to Image

This task involves writing an array of numbers into an image while taking care of it being grayscale or colored. Again, OpenCV library has been used for the major part of this task.

```
def ImageWrite(src, Target_Dir: str, image_file_name: str, is_rgb: bool) → None:
    """
    Write an array into image-file and return the filepath.\\
    Note: Put is_rgb as False only if the image is Grayscale.

    Arguments:
    src = input image array
    Target_Dir = Path to the target directory
    image_file_name = name of the file the image needs to be written into
    is_rgb = a boolean value denoting whether the image is rgb or not;
    |         if not, image will be considered grayscale
    """

    assert os.path.exists(Target_Dir), f"{Target_Dir} is invalid!"
    image_path = os.path.join(Target_Dir, image_file_name)

    if (is_rgb):
        | src = Convert(src, is_BGR = False, to_RGB = False, is_RGB = True, to_BGR = True)
    cv2.imwrite(image_path, src)
    return image_path
```

Figure 2: Function for writing to an image from an array

The inputs to the function are the source array which needs to be written to the image, target directory and filename along with a boolean value denoting whether the image should be RGB (if False, then the image is taken to be grayscale - this helps in storing the image in BGR/Grayscale format). The function returns the local path to the location of the image. Outputs can be seen in the main jupyter notebook.

### 0.3 Task 3 - Brightness Change

Given an image, the aim of this task is to change its brightness by some factor which will be added uniformly across the image. The code is designed to make this factor - *b\_factor* - user defined.

$$R(x, y) = S(x, y) + bf$$

On setting various values for the brightness factor, we get a range of images with varying brightness.



Figure 3: Grid of images depicting varying brightness where bf = brightness factor

The shift is evident in the histograms. (For the histograms with bf = -100 and bf = +100, the y-range is very large and to fit it alongside other ones, the bars are not very clearly visible).

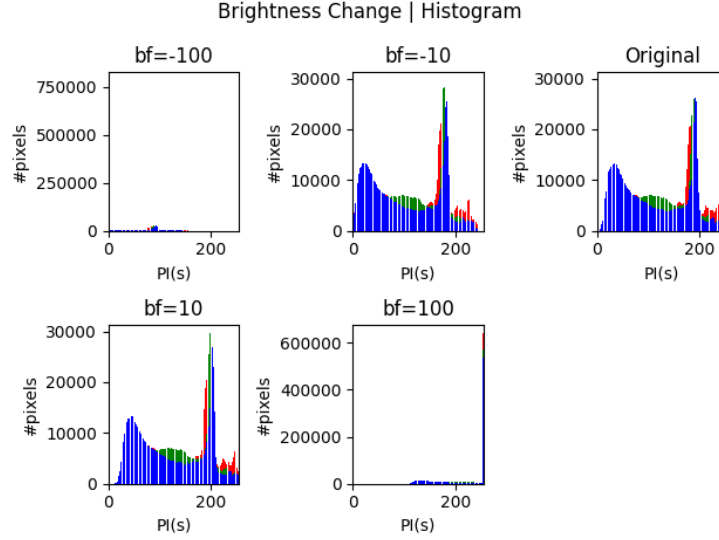


Figure 4: Grid of histograms depicting effects of varying brightness where  $bf$  = brightness factor and  $PI(s)$  = pixel intensities

Code for this task has not been attached in the report but can be found along with the plots in the main code file.

**Note:** For this task, since we are adding some factor to the image to increase or decrease its brightness, the pixel values tend to get out of the range which makes plotting and visualising the image very difficult. To counter this, while adding, the pixel values are clipped to stay within range  $[0, 255]$ . Finally, the datatype of the array is converted to `uint8` for visualisation purposes.

## 0.4 Task 4 - Contrast Stretching

Contrast stretching is one method to beautify an image to be able to draw as much information as possible from the image. This technique is applied on the images which are of low-contrast. Low-contrast implies that the range of values that the pixels of the image takes is limited to a very small region between 0 and 255; in effect the histogram looks very skewed or shrunk to a small region. Since all the color intensities are very near to one another the perceivable difference in the image is very less, hence low-contrast.

The aim of this task is to convert a low-contrast image into a high-contrast one. Two methods to do it:

- **Linear:** Involves simply multiplying the image by a constant and then clipping for the values that go out of the permissible range.

$$R(x, y) = \alpha S(x, y)$$

- **Non-Linear:** Involves piece-wise stretching of the pixel values. Here, for each short range (three) of permissible values (0 to 255) a separate function is defined for performing the stretching.




Figure 5: Formulation of piece-wise contrast stretching - taken from Digital Image Processing textbook

The code I wrote gives the user to choose any of the above methods to do contrast stretching. It also, infact allows the user to perform **contrast compression** (reverse of contrast stretching). The results I show are of an image which has been converted to a low-contrast one and then changed back to stretched version using contrast stretching. The image for non-linear method is on the next page (Figure.7).

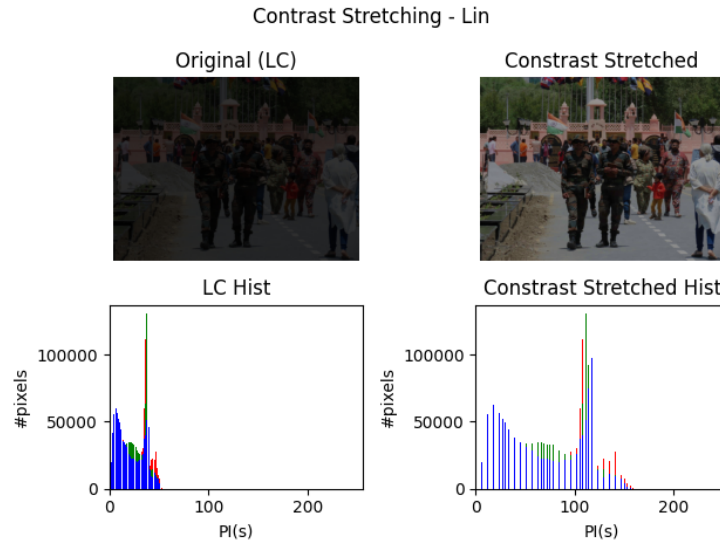


Figure 6: Contrast Stretching using Linear method

## 0.5 Task 5 - Color to Grayscale

This task involves converting the image space of an image - taking it from RGB space to Grayscale space. One major change that is seen here, and which is quite natural, is that the number of channels decrease from 3 (in case of RGB) to 1 (in Grayscale). This also implies that the grayscale image occupies less memory than its coloured counterpart. Now, although the number of channels has decreased, the range of values that the pixels can take still remains the same (i.e., 0 to 255). A few methods to do this:

- **Luminance-based:** Uses the colorimetry principles and the considers biological sensitivity of the human eye in perceiving R, G and B channels. This method preserves the luminance aspect of the image (light directions, etc.) and therefore looks natural. The weights associated with R, G and B channels here are 0.2126, 0.7152 and 0.0722 respectively. (ref.: Wikipedia).
- **Arithmetic Mean-based:** A very straightforward method taking mean of all the three channel pixel values and that mean becomes the grayscale pixel value.
- **Max-of-all -based:** Here we take the maximum of the three channel values for each pixel in the

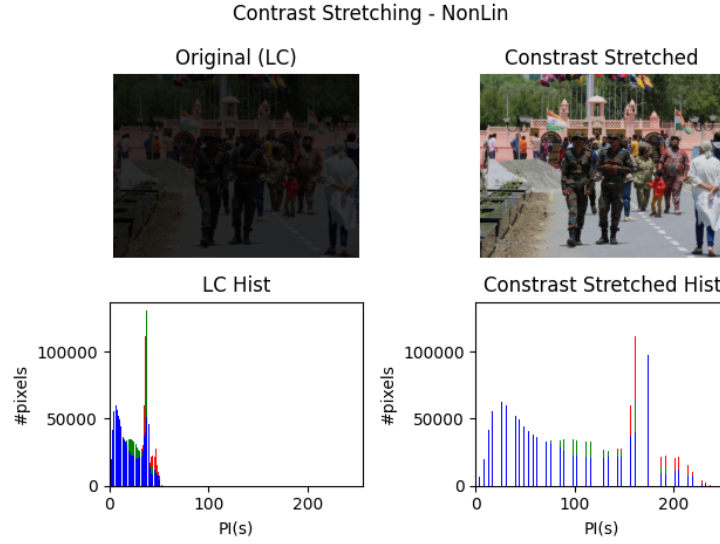


Figure 7: Contrast Stretching using Non-Linear method

color image and this value then becomes the resultant gray value. Renders an image more towards the brighter end of the spectrum.

- **Min-of-all -based:** Similar to max-of-all method, here we take the minimum. Renders an image more towards the darker end of the spectrum.
- **Desaturation-based:** Here we use the max-of-all and min-of-all methods to find the arithmetic mean of both and give that value to our grayscale image.
- **Weighted Average-based:** Here we put any arbitrary weights to R, G and B channels and see the effects (I have made this user-defined inputs).

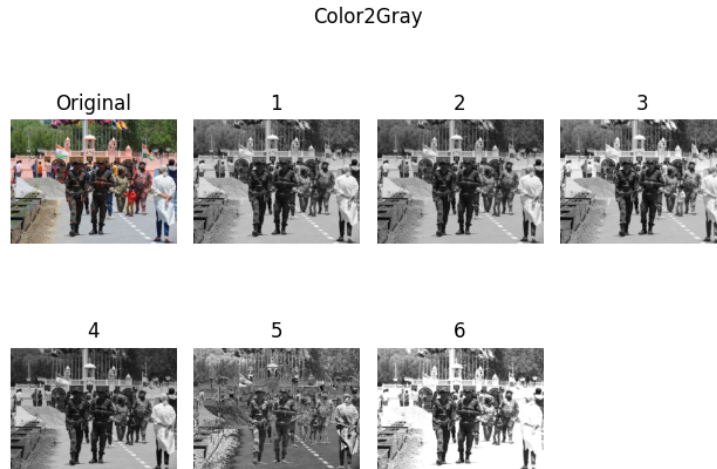


Figure 8: Color to Grayscale using the six methods listed above (in the same order)

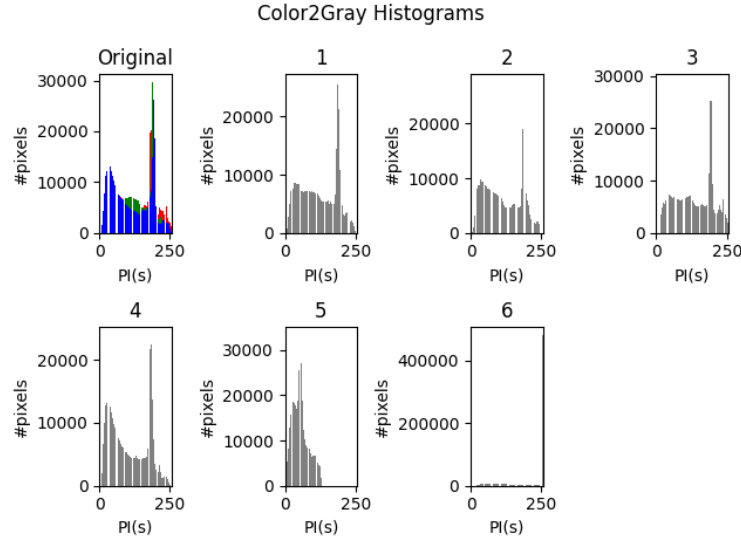


Figure 9: Histograms for color to grayscale using the six methods listed above (in the same order)

## 0.6 Task 6 - Grayscale to Color

This was a tricky task but interesting because there was no right answer to it (ofcourse except the actual image) and thus, a lot of space for experimentation was there. I made use of, as was mentioned, pseudo-coloring to generate colors for a grayscale image. The method I thought of was quite straight forward and involved making gradients for the three primary colors (R, G and B in our case) using fading techniques. For instance, I started off with Blue, then while it was gradually set to fade, I was proportionally increasing Green. Then for the next part of the gradient, I started to fade Green and gradually pushed Red into picture finally culminating in a dark or bright tone. I generated the following gradients:

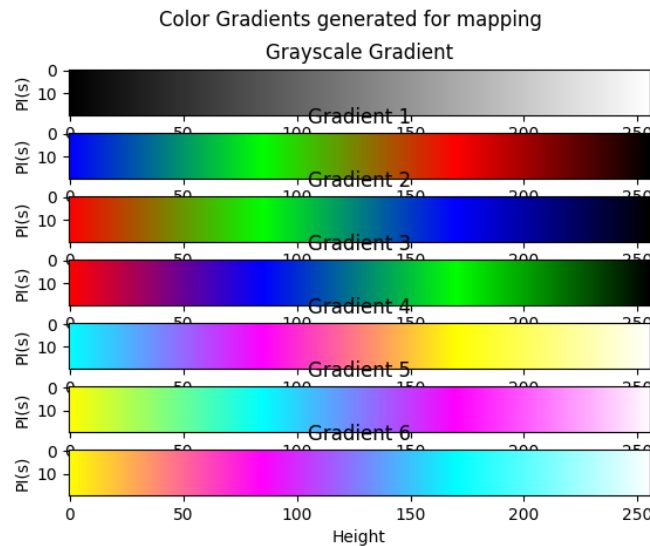


Figure 10: Color Gradients generated for pseudo-color mapping

The gradient-wise images are as follows:  
And some extras...

### Gradient-wise Color Images




Figure 11: Color Gradients-based color Images

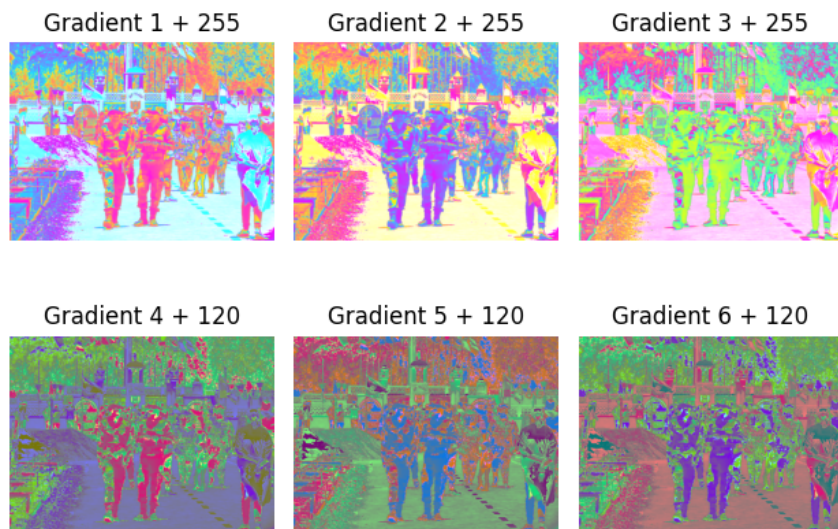


Figure 12: Some extra images obtained after playing around with the ones in Figure.11



**Note:** Clipping was required in all the above generated color images.

## 0.7 Task 7 - Chroma Keying

Chroma keying is the name given to the process of removing greenscreen background from an image and replacing it with any other background that one wants. Therefore, this at least needs 2 images to start with - one foreground (with greenscreen) and the background (which will replace the greenscreen).

It is quite clear that to remove the background, we first need to know the numeric distribution of the pixel intensities in both the images, and what better than a histogram to do it. I used the following images:

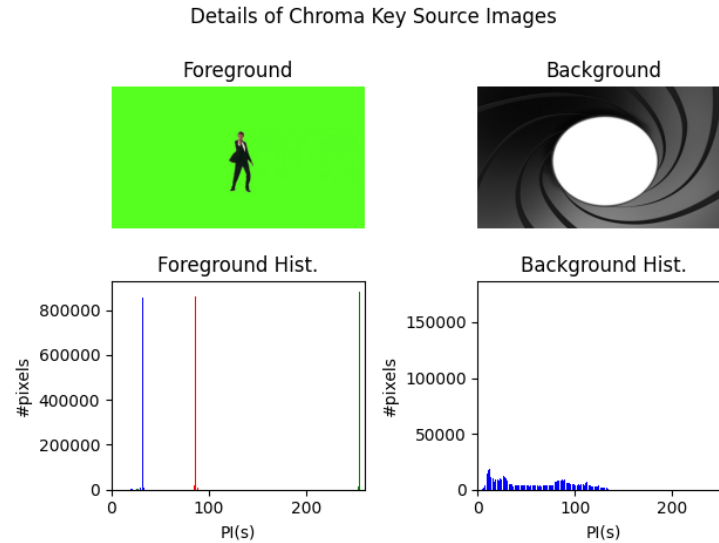


Figure 13: Chroma Keying input images

As we can see in the images, the green pixel count of foreground image is quite high and almost completely towards the lighter end of the spectrum. In my method, I conveniently removed the pixels with values above 210, and superimposed the background over it to get the required effects. The method looks as follows:

```
image = src_fgd.copy()
src_fgd = np.where(src_fgd[:, :, 1] > 210, 255, src_fgd[:, :, 1])
image[src_fgd == 255] = src_bkg[src_fgd == 255]
```

Figure 14: Chroma Keying algorithm

The algorithm commented out as *Aliter* also does the same work but is implemented using **for-loops** whereas, the one which isn't commented out is the vectorised version of the same thing.

On doing this, I get the final image which looks as shown below:

**Note:** The function has the ability to take care of images with unequal shapes (H and W values). For this the function makes use of the **Resize** functionality of OpenCV (a one line code).

## 0.8 Task 8 - Reading from and Writing to a Video

Videos are essentially nothing but a sequence of frames (images) being rapidly changed by sampling time. The rate at which the frames move (FPS) and the duration of the video determine the number of frames that the video includes. Taking into consideration the persistence of vision which is around  $\frac{1}{16}^{th}$  of a second, that is, 16 frames per second, any FPS above that should be fine.



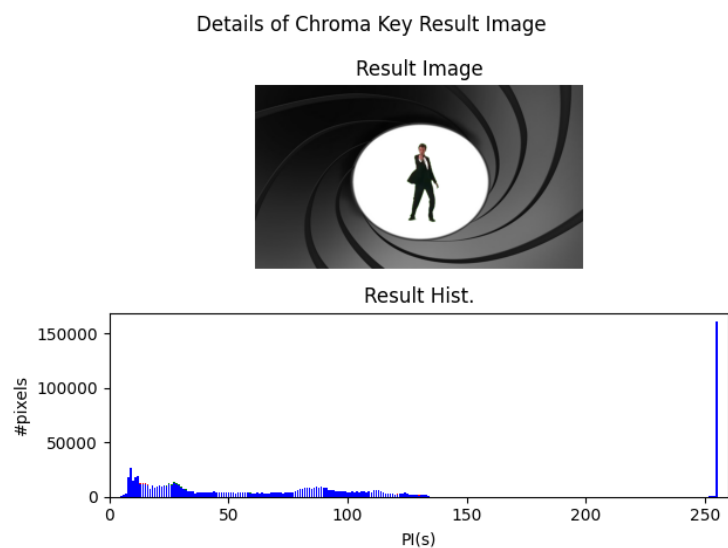


Figure 15: Chroma Keying output image

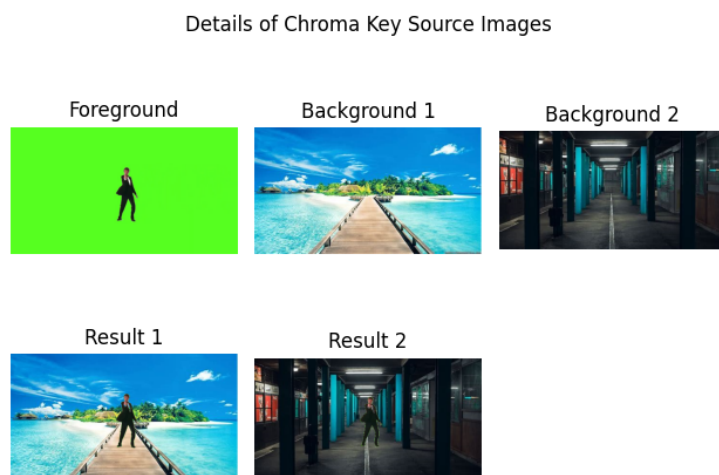


Figure 16: Some other trials with the same foreground but different *suitable* backgrounds

To carry this task out, I made use of OpenCV functions that help one in reading and writing into images; specifically, I used **VideoCapture** and **VideoWriter** to do the reading from and writing to tasks respectively.

While reading the video, I stored the images - as arrays - in a concatenated form as a list. This was for three reasons - one, to not occupy memory in the machine, second, to sustain the sequence and third, for easier access across various other functions (such as writing or editing). While writing to the video, I accessed the stored frames - as arrays - in a sequential manner. I have kept FPS as well as the codec as a user-defined variable for one to play around with (in case of FPS) and to select the right video extension (in case of codec).

**Learning:** There was one major learning worth mentioning separately while writing to a video which was to realise that OpenCV understands the image in WHC (weight, height and channels) format when it comes to giving these params as inputs. It took me a while to figure out this since I was giving the frames in HWC format which kept on shooting errors. The original video is [here](#) and the link to the recreated video can be found [here](#).

**Note:** The input video file has to contain all frames in colored format.

## 0.9 Task 9 - Creating a Special-Effects Video

This was one the most interesting tasks which involved using some special effect such as fading, sliding or any other had to be used to make a video wherein one image makes a transition into another. I wrote a function facilitating one to play around with either of the effects (with custom duration and FPS) using the following techniques:

- **Fade:** Using the input FPS and duration, I calculate the number of frames that the video will have and then iterate over that number to generate those frames using the formula:

$$R(n) = (1 - \frac{n}{N})S_1 + (\frac{n}{N})S_2$$

where  $S_1$  and  $S_2$  are source images,  $N$  is the total number of frames and  $n$  is the iteration variable which goes from 0 to  $N-1$ .

- **Slide:** Here I did not make use of FPS in determining the number of frames as by choosing this effect the function automatically selects the total number of frames to be equal to the columns of one of the images. This is because the driving logic for this effect is to copy the new image column-by-column over the first image thereby creating a *sliding* effect.

Outputs videos can be found here:

- [Fade](#)
- [Slide](#)

Few things to **note:** The shape of both the images has to be equal otherwise these effects will not work. But, the user does not need to worry about this! The function for this task will take care of the unequal shapes by detecting it and employing the **Resize** functionality of OpenCV. One thing that the one does need to take care about while putting in inputs is that the function will not accept 2 different *types* of images - color and grayscale; either both have to be color or both need to be grayscale.

## 0.10 Misc. Functions

I created some miscellaneous or helper functions that make some tasks easier:

- **Generate Histograms:** I wrote a function that takes an image and its type as an input and generates appropriate histogram for it. Simply calling this function during visualisation saves some repetitive lines of code!



Figure 17: Input images to Task 9

- **Convert:** I wrote this function which aims to convert the type of image from BGR to RGB and vice-versa. Although only 3 lines of code, this comes a lot in handy when we realise that to visualise an image using Matplotlib which has been read using OpenCV, conversion is very necessary. This is because OpenCV reads an image in BGR format and Matplotlib visualises it using RGB format, thereby rendering a very weird image as an output.

Finally, as a **note**, all the task has been majorly carried out using just Numpy and Matplotlib. OpenCV has only been used for tasks such as reading, writing or resizing.